






ali 
608 posts



ali 

@waterloo_intern

ml research, kernels, and the occasional peer-reviewed shitpost
inference [@baseten](#) || eng [@uwaterloo](#)

 San Francisco  [github.com/AliesTaha](#)  Joined October 2024

108 Following 26.3K Followers

Mention

Follow

Posts

Replies

Articles


Media



Pinned



ali 
@waterloo_intern · Jun 26



Article

some notes on writing the fastest video kernel in the world

in this worklog, I explain how we retrofitted sparsity into a model at inference time, yielding the world's fastest video generation kernel. we iteratively optimize a frontier oss model by 54x the...

Don't miss what's happening

People on X are the first to know.

Log in


Sign up

I spent 48 hours with the Kimi K3 modeling code.

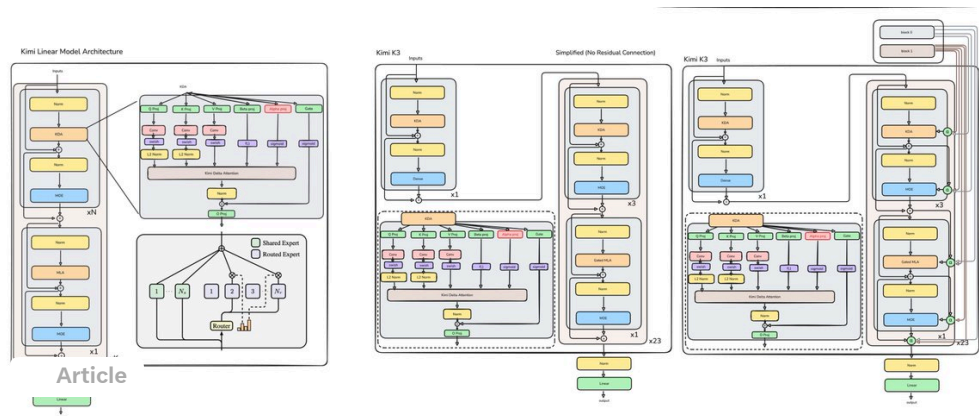
It took:

- 650 mg of caffeine (mandatory)
- 40 cans of LaCroix (optional... world record ??)
- 8 papers
- 6 months off my lifespan

Finally grokked the entire lineage of Kimi K3 and how we got here... every single step, since [Show more](#)



 ali  @waterloo_intern · Jul 27



22580: From GPT2 to Kimi3, Explained

Twenty-two thousand five hundred and eighty. That's how many GPT-2 (2019) models fit inside KimiK3 (2026). We scaled up by a factor of 22,580 in seven years. But is it just... scale? In this worklog,...

 221

 638

 8.7K

 1.9M




ali  @waterloo_intern · Jul 27

Don't miss what's happening

People on X are the first to know.

Log in

Sign up



22580: From GPT2 to Kimi3, Explained

Twenty-two thousand five hundred and eighty. That's how many GPT-2 (2019) models fit inside KimiK3 (2026). We scaled up by a factor of 22,580 in seven years. But is it just... scale? In this worklog,...



ali 191

@waterloo_intern · Jul 16

1.9K

11K

4.4M

it took 3 (and a half / intern) engineers 12 weeks of work to get:

- > the fastest video inference engine in the WORLD
- > the most guardrailed video model in the WORLD

i'll give you 1 BILLION TOKENS to go break it.

we generate a 5 second video in sub 2.5 seconds
every [Show more](#)

32

2

288

73K



ali @waterloo_intern · Jul 16

reading this paper changed my priors on continual learning (that's the closest you're getting to an admission of guilt)

i mean, the problem itself seems quite simple

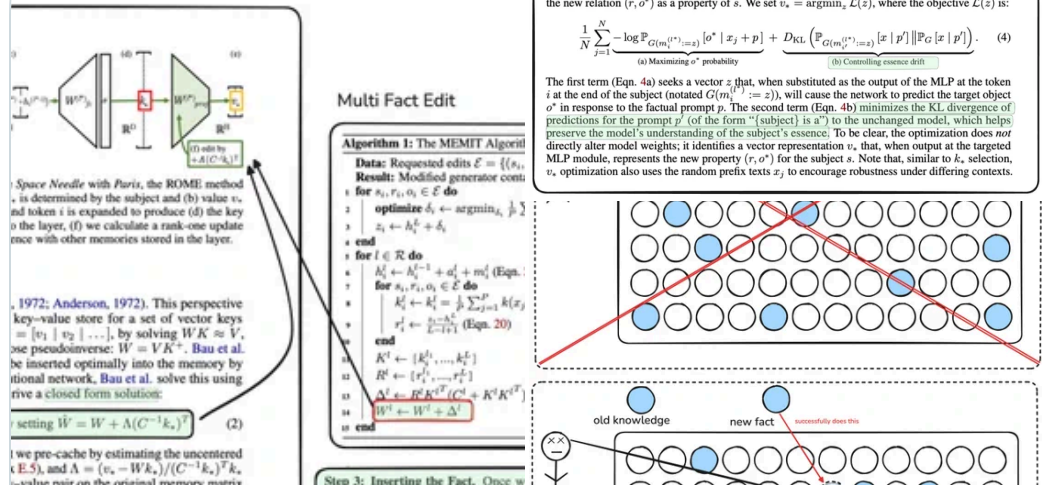
- > person uses model
- > model learns and has a couple of 'aha' moments
- > push those new facts into the weights [Show more](#)

Don't miss what's happening

People on X are the first to know.

Log in

Sign up



See @waterloo_intern's full profile

Sign up

Log in

20

46

724

196K

1

Don't miss what's happening

People on X are the first to know.

Log in

Sign up