


 205987.2 TFLOPS

Hugo Larcher

hllarcher



3



22



10

Follow

 157 followers · 62 following

 hugoch  hugoch  hllarcher

 hllarcher.bsky.social

AI & ML interests

None yet

Recent Activity



Updated a Space 1 day ago

[huggingface/anatomy-of-frontier-...](#)


Published a Space 1 day ago

[huggingface/anatomy-of-frontier-...](#)


Published an article 2 days ago

[Anatomy of a Frontier Lab Agent Intrusion...](#)

View all activity

Organizations



Buckets 2



↑↓ Sort: Recently updated



hllarcher/hf-mount-bundle-test

5.41 GB



hllarcher/mmap-repro-test

15.2 GB

Posts 2



Post

500



GH200 cooking time 🧑🍳🔥!

[...read more](#)

View all Posts

Articles 4



Article

90



Anatomy of a Frontier Lab Agent Intrusion: A Technical Timeline of the July 2026 Incident

View all Articles

Papers 2

SmolLM2: Redefining small and efficient multimodal models



SmolLM2: When Smol Goes Big — Data-Centric Training of a Small Language Model

Ben Allal¹, Anton Lozhkov², Elie Bakouch³, Gabriel Martín Blázquez⁴, Guilherme P. Tunstall⁵, Andrés Marafioti⁶, Hynek Kydlíček⁷, Agustín Piqueres Lajarin⁸, Vaibhav Srivastava⁹, Caleb Fahlgren¹⁰, Xuan-Son Nguyen¹¹, Clémentine Fourier¹², Ben Burtenshaw¹³, Hugo Larcher¹⁴, Haojun Zhao¹⁵, Cyril Zakka¹⁶, Mathieu Morlon¹⁷, Colin Raffel¹⁸, Leandro von Werra¹⁹, Thomas Wolf²⁰

[//hf.co/collections/HuggingFaceTB/smolLM2-6723884218b0da64b34d](#)

Abstract

Large language models have facilitated gains in many applications of artificial intelligence, but their inherent largeness makes them difficult to deploy in resource-constrained settings. In this paper, we present the SmolLM2 series of models, which are, by their nature, *large*, in the sense that they have many parameters (more than ~10 billion), but they are also *small*, in the sense that they are trained on a smaller dataset and are more efficient to run. This results in a model that is both powerful and efficient, making it a good choice for inference in resource-constrained environments.

arxiv:2502.02737

Running on **CPU UPGRADE**

Patrouille de France Tracker ➦

Jan 30

Running

Cloud provider bandwidth 🚚

Explore cloud bandwidth benchmarks for model downloads

Apr 11, 2025

Sleeping

Tgi Bench 🖥️

Apr 2, 2025

Build error Agents

Hello World 💬

Dec 11, 2024

 hlarcher/xet-bench2

Updated May 29

 hlarcher/xet-bench

Updated Nov 5, 2025

 hlarcher/opt-125m-lfs

 Text Generation • Updated Feb 21, 2025 • ⬇ 5

 hlarcher/opt-125m

 Text Generation • Updated Feb 21, 2025 • ⬇ 4

 hlarcher/falcon-7b-v100s

 Text Generation • Updated Jul 25, 2023 •  6

 **Datasets** 7 

↑↓ Sort: Recently updated

 hlarcher/random-data-2

Updated May 27 •  69

 hlarcher/random-data

Updated May 27 •  84

 hlarcher/cloud-provider-bandwidth

 Viewer • Updated Apr 11, 2025 •  148 •  16

 hlarcher/inference-benchmark-xet

 Viewer • Updated Feb 14, 2025 •  120k •  47

 hlarcher/inference-benchmark

 Viewer • Updated Jan 8, 2025 •  120k •  622 •  1

 hlarcher/thisisatest

 Viewer • Updated Oct 29, 2024 •  1 •  22

 hlarcher/share_gpt_small

 Preview • Updated Oct 14, 2024 •  27

 System theme

TOS

Privacy

About

Careers



Models

Datasets

Spaces

Pricing

Docs